

Deployment of a Microservices Architecture based Web App by using Docker, Kubernetes, and Azure DevOps

Anik Saha, Mursheda Baby, Shalim Sadman
High Integrity System (MSc.)
Frankfurt University of Applied Sciences

Abstract:

Microservices are a method of building cloud applications that is both architectural and organizational. Microservices break down an application into smaller independent pieces that need a way to be managed. Containerization with Kubernetes orchestration and management is designed to support microservices. Microsoft Azure offers Azure Kubernetes Service that simplifies managed Kubernetes cluster deployment in the public cloud environment and manages the health and monitoring of managed Kubernetes service. This project includes three servers that leverage container-based technologies inside a microservice architecture, making resource utilization easier and more efficient. The idea is to explain containerized application deployment into the Kubernetes cluster with Azure DevOps, Dockerhub, and a Github repository.

1. Introduction

Microservices are based on the concept of deploying individual apps, each of which delivers a specific service and then connecting these microservices to form a larger service. Microservices make it easier to deploy and upgrade individual apps, making it easier to create sophisticated services [16]. The usage of Docker to host microservices is becoming increasingly popular. Docker provides a command-line interface (CLI) for creating Docker images, which may then be used as templates for creating containers [8]. Images can be maintained locally or published to the Docker Hub, where they can be made privately or publicly visible. It's simple to install and configure one or more Docker containers on a host to establish an integrated service. However, managing a cluster of compute nodes, each of which hosts several containers, is more complicated, and some ways of cluster management are required. Kubernetes is a container orchestration system that may be used in various contexts, including Google Compute Engine and Microsoft Azure.

The Azure Kubernetes Service (AKS) is the most convenient method to start with Kubernetes on Azure. It is unquestionably an ideal platform for developers to build modern apps using Kubernetes on the Azure architecture, with Azure Container Instances being an excellent alternative for public cloud container deployment. Azure Container Instances allow developers to install and execute their applications on Kubernetes architecture with less stress.

Organizations are increasingly modernizing application development by incorporating open source technologies into a comprehensive architecture for cloud delivery of high-quality workloads [4]. This document will introduce some basic Kubernetes ideas and provide step-by-step guidance on deploying a simple web application using Azure DevOps with Kubernetes and Docker.

2. System overview

This section contains a brief summary of the technology, processes, and approaches used in the system. Finally, the techniques and architecture utilized to construct the overall system are described in depth.

2.1 Microservices Technology and Architecture

Microservices are mostly utilized to alleviate issues caused by monolithic systems. Microservices is a design approach that reflects an application's structure by merging all of the application's autonomous services. Microservices architecture allows the distribution of an application among services that are independent of one another. The following are some key characteristics of a microservice:

High availability – Every microservice should respond to all requests in a reasonable amount of time.

Partition tolerance – If a microservice instance fails, the system should not be affected. Should ensure that the availability of microservices is consistent.

Eventual consistency – When data in a single microservice changes, it should eventually be propagated to other microservices that are relevant.

Asynchronous communication -- Microservices should not connect asynchronously. Instead, they should delegate communication to a message broker [15].

2.2 Communication Between Microservices on Kubernetes

There are several approaches to expose a Kubernetes-based application:

The most common technique is to use a Kubernetes service, a network abstraction or logical entity of pods. There are a variety of services available, each with its own set of capabilities; this project makes use of the following:

ClusterIP uses Kubernetes' default service option and exposes the service on a cluster-internal IP.

NodePort exposes the application to the outside world via a node-level static IP and port combination.

LoadBalancer makes the application available as a service to a cloud-based load balancer [17].

2.3 Kubernetes Basics:

Kubernetes is an open-source framework that helps to manage containerized workloads and services, including declarative and automated setup capabilities [2]. Kubernetes consist of numerous components that are completely unaware of each other. All of the components interact with one another via the API server. Before exposing metrics that can be gathered for further monitoring, each component has a specific purpose [10].

Let's go over the Kubernetes fundamentals that were used to create this project:

A "pod" in Kubernetes is a collection of functionally connected containers.

A "service" is a collection of connected pods that perform the same set of tasks. Kubernetes gives each Pod its IP address and a collection of Pods a single DNS name. IP addresses for pods are not stable because each new pod is allocated a new IP address; hence, a direct connection between pods is not often viable. On the other hand, services have their relatively constant IP addresses; consequently, an external resource requests a service rather than a pod, and the service forwards the request to an available pod [5].

For non-confidential data, a ConfigMap is a key-value pair storage API object. Pods can use ConfigMaps as environment variables, command-line options, and disk configuration files. With a ConfigMap, we can isolate environment-specific configuration from container images, making the apps more portable.

A Kubernetes deployment is a resource object that allows apps to get declarative updates in Kubernetes. A deployment allows defining an application's life cycle, including which images to use, how many pods should be present, and how they should be updated.

The PersistentVolume subsystem provides a user-friendly API that separates storage provisioning from storage consumption. Two new API resources that help with this are PersistentVolume and PersistentVolumeClaim.

A secret is a small amount of confidential data, such as a password, token, or key. Alternatively, such information might be included in a Pod specification or an image. Secrets can be created by users, and the system will also generate certain Secrets.

When a workload demands surges or falls, Kubernetes can automatically raise or reduce the number of pod replicas providing the workload, known as horizontal scaling. This is a dynamic feature with a reconciliation loop that uses experimental measurements to drive the workload's capacity toward the capacity established by the workload's owner.

2.4 System Architecture

Azure DevOps, which manages container deployment to the Kubernetes cluster, serves as the system architecture's backbone. As a source control repository, we used Github, and as a content repository, we used DockerHub. The Azure DevOps pipeline is in charge of container creation and deployment. When the developer updates the Github source, the pipeline is started. To create apps and communicate with docker hub, GitHub, and Azure Kubernetes services, Azure DevOps uses plugins. The services are dockerized and hosted in a public DockerHub repository before being deployed to an Azure Kubernetes cluster. A load balancer service in a Kubernetes cluster directs traffic to available pods. Figure 1 depicts our high-level system design.

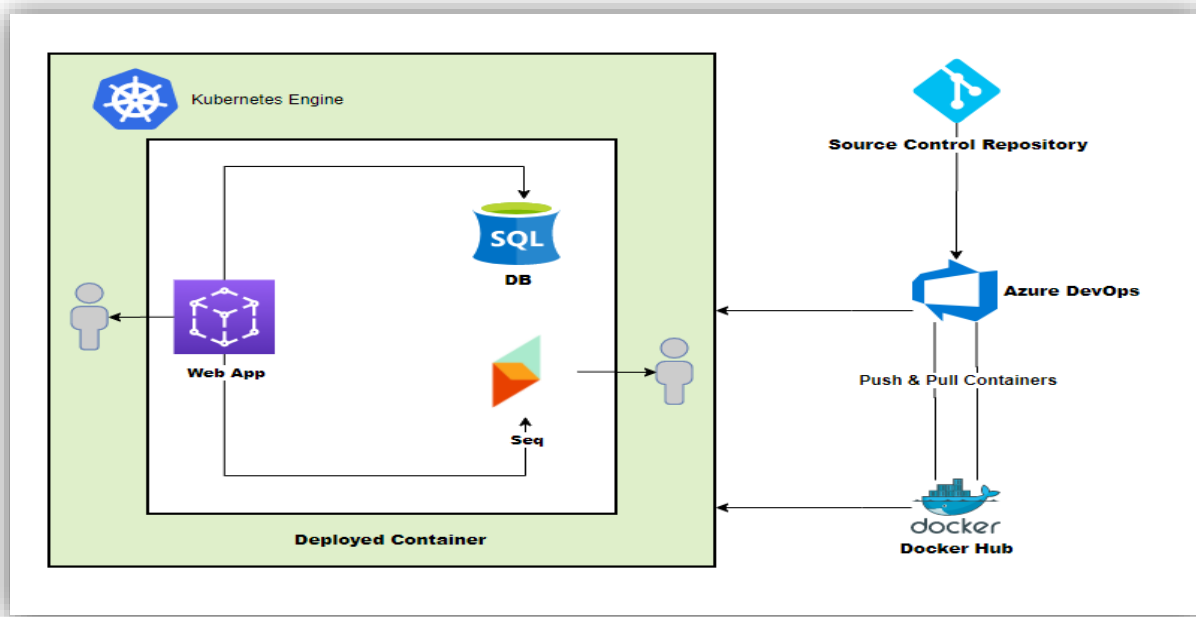


Figure 01: System Architecture

2.5 Tools and Methodologies

Deployment tool: Kubernetes, Azure DevOps

Programming Language: C#, HTML, CSS, JavaScript

Source Control Repository: GitHub

Container repository: Docker Hub

Log platform: Seq

Testing tool: k6

3. Implementation:

The section details the implementation steps towards achieving our desired objectives. Three servers have been used for the implementation. They are Web app (two pods), log server (1 pod), and MySQL server (1 pod). Cmd is used to run the Kubernetes command.

3.1 About MVC Web Application:

It's a simple web application. There are only three pages in this section. The first is the homepage. The Privacy page is the second, and the Product Information System is the third. We just save product names in our system because our whole focus is on Kubernetes right now. A product's name can be easily added, deleted, or edited by anyone.

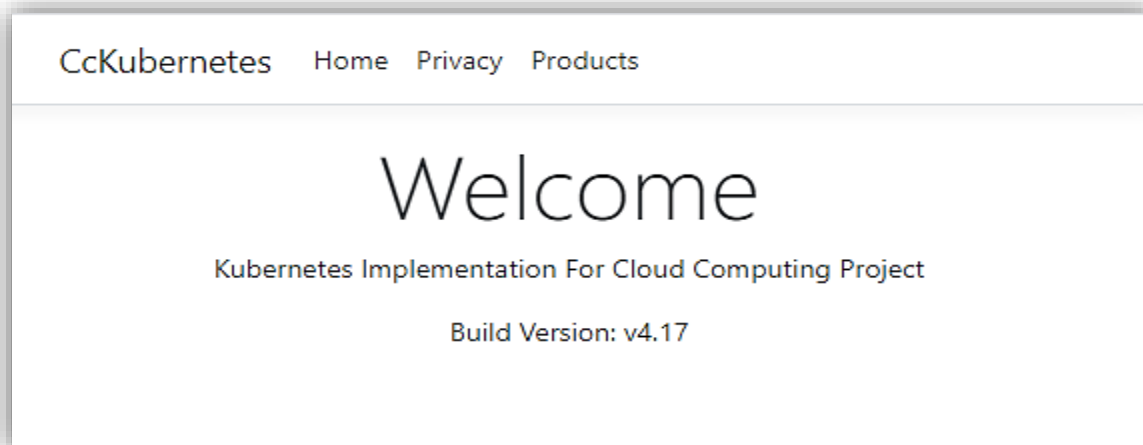


Figure 02: Home Page




Figure 03: Privacy page

The code of the home and privacy page is available [here](#).

Then we are adding create, delete, edit of product.

The screenshot shows the 'Create Product' form in the CcKubernetes application. The header includes the application name 'CcKubernetes' and navigation links for 'Home', 'Privacy', and 'Products'. The main heading is 'Create Product'. Below it is a form with a 'Name' label and an empty text input field. A blue 'Create' button is positioned below the input field, and a blue link 'Back to List' is at the bottom left.

CcKubernetes Home Privacy Products

Create Product

Name

Create

[Back to List](#)

Figure 04: Creating product

The screenshot shows the 'Edit Product' form in the CcKubernetes application. The header includes the application name 'CcKubernetes' and navigation links for 'Home', 'Privacy', and 'Products'. The main heading is 'Edit Product'. Below it is a form with a 'Name' label and a text input field containing the value 'test'. A blue 'Save' button is positioned below the input field, and a blue link 'Back to List' is at the bottom left.

CcKubernetes Home Privacy Products

Edit Product

Name

Save

[Back to List](#)

Figure 05: Editing product

The screenshot shows the 'Index' page in the CcKubernetes application. The header includes the application name 'CcKubernetes' and navigation links for 'Home', 'Privacy', and 'Products'. Below the header is a link 'Create New'. The main content is a table with one row. The table has a header row with 'Name' and a data row with 'test' and a set of links 'Edit | Details | Delete'.

CcKubernetes Home Privacy Products

Index

[Create New](#)

Name
test

[Edit](#) | [Details](#) | [Delete](#)

Figure 06: Added product

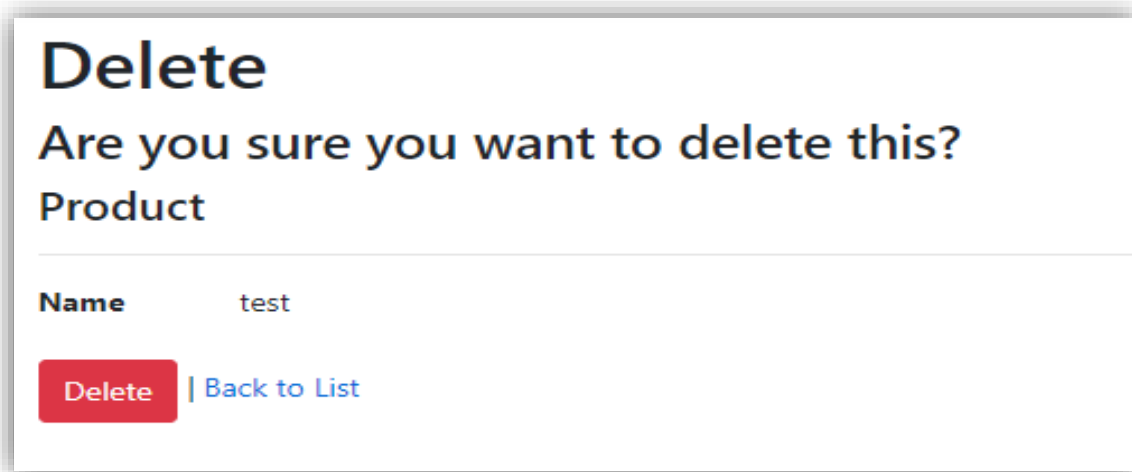


Figure 07: Deleting product

The code of the whole product information system is [here](#).

3.2 Creating a Docker Image (Creating the Docker containers for the app)

The following set of commands inside 'Dockerfile' is used to build and run the web app into the [Docker container](#).

```
FROM mcr.microsoft.com/dotnet/aspnet:5.0-buster-slim AS base
WORKDIR /app
EXPOSE 80
EXPOSE 443

FROM mcr.microsoft.com/dotnet/sdk:5.0-buster-slim AS build
WORKDIR /src
COPY "CcKubernetes.csproj" .
RUN dotnet restore "CcKubernetes.csproj"
COPY . .
RUN dotnet build . -c Release -o /app/build

#RUN apt-get install curl

FROM build AS publish
RUN dotnet publish "CcKubernetes.csproj" -c Release -o /app/publish

FROM base AS final
WORKDIR /app
COPY --from=publish /app/publish.
ENTRYPOINT ["dotnet", "CcKubernetes.dll"]
```

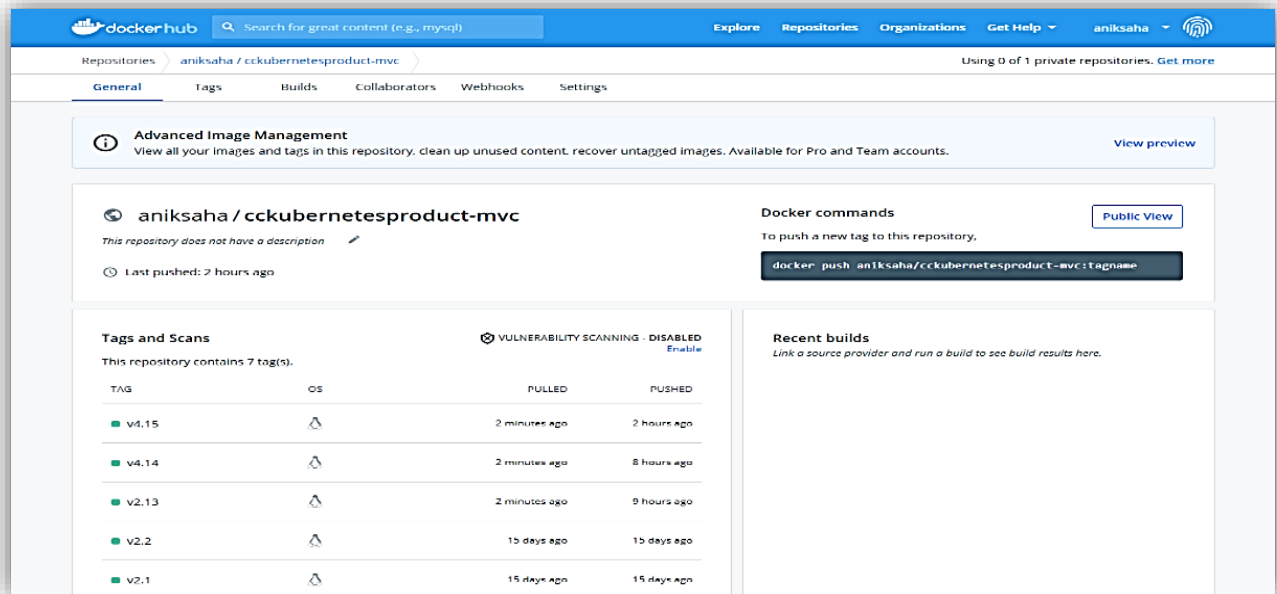


Figure 08: Image on Docker Hub

For SQL server and Seq log server, Official images for [Microsoft SQL Server](#) and [Seq log server](#) have been used respectively for Docker Engine.

Log server - Seq:

Seq is a modern structured logging system that includes message templates. It combines free-text and regular expression searches with intuitive expression-based filtering. Instead of wasting time and effort trying to extract data from plain-text logs using inaccurate log parsing, the properties of each log event are collected and sent to Seq in a clean JSON format. Message templates are natively supported by ASP.NET Core, and because our application is built on it, it provides the optimum diagnostic logging for our platform [9].

It is a third-party log for this project. We extract this from the docker image. Two ports are set in this seq log to get logs that operate through the web app connection. Seq provides the visibility to quickly detect and fix problems in complex systems and microservices.

These logs are delivered over the network to Seq, who displays and searches them:

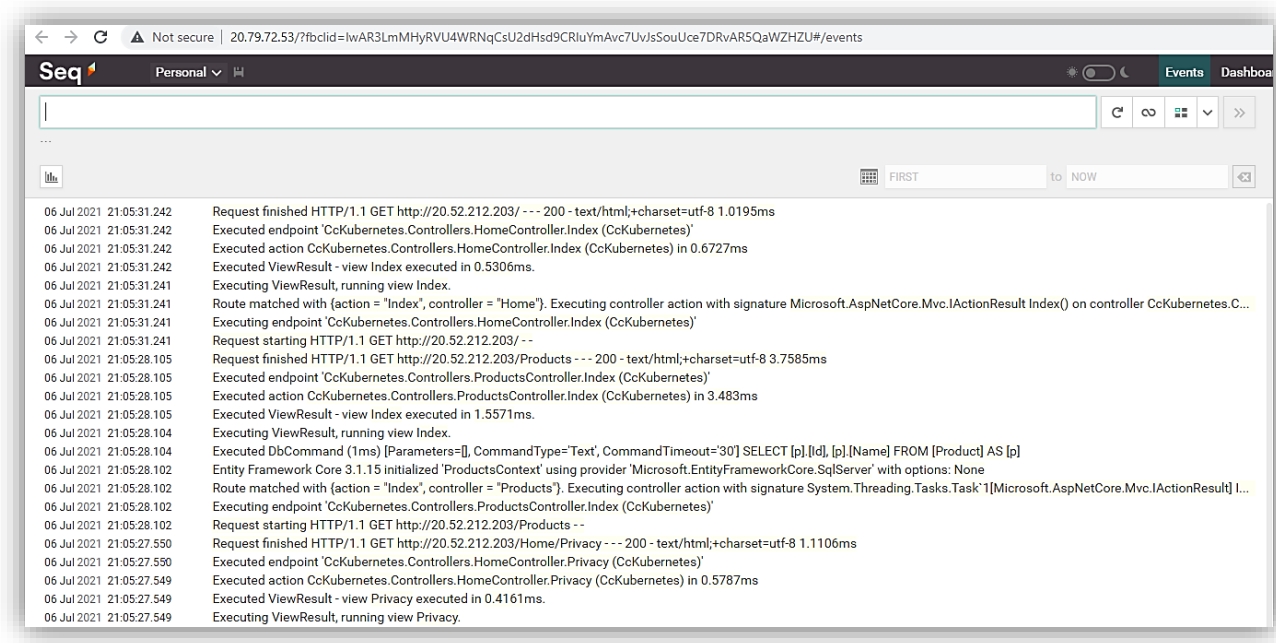


Figure 09: Log query of web application

Here we have given the code of connecting code between seq and our mvc app.

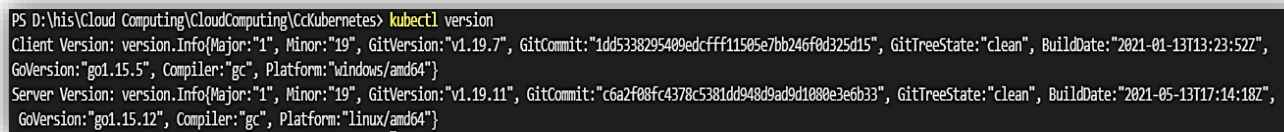
```
public static IHostBuilder CreateHostBuilder(string[] args) =>
{
    Host.CreateDefaultBuilder(args)
        .UseSerilog((ctx, provider, loggerConfig) =>
        {
            loggerConfig
                .ReadFrom.Configuration(ctx.Configuration) // minimum levels defined per project in json files
                .Enrich.FromLogContext()
                .WriteTo.Console()
                .WriteTo.Seq($"http://{ctx.Configuration.GetConnectionString("Seq")}:{ctx.Configuration.GetCo
nnectionString("SeqPort")}");
        })
        .ConfigureWebHostDefaults(webBuilder =>
        {
            webBuilder.UseStartup<Startup>();
        });
}
```

3.3 Kubernetes on Docker Desktop

Kubernetes is a container-based system. As a result, an account on the Docker hub was first created to push the microservices image and MySQL image to the Docker hub, which is required for Kubernetes deployment. We aim to build up our Kubernetes in a local workstation before playing with the Azure Kubernetes service. It's also simple to set up Kubernetes with Docker Desktop to build and push docker images into the docker hub. We deployed the Kubernetes cluster during Docker installation, which runs all Kubernetes components in containers.

Docker Desktop will download all of the Kubernetes images and start things up in the background. When it's ready, two green lights will appear at the bottom of the settings screen, indicating that Docker and Kubernetes are running.

The Kubectl is then verified with the following command:

A terminal window showing the command 'kubectl version' and its output. The output displays both client and server version information, including major, minor, and patch versions, as well as Git commit hashes and build dates.

```
PS D:\his\Cloud Computing\CloudComputing\Cckubernetes> kubectl version
Client Version: version.Info{Major:"1", Minor:"19", GitVersion:"v1.19.7", GitCommit:"1dd5338295409edcfff11505e7bb246f8d325d15", GitTreeState:"clean", BuildDate:"2021-01-13T13:23:52Z",
GoVersion:"go1.15.5", Compiler:"gc", Platform:"windows/amd64"}
Server Version: version.Info{Major:"1", Minor:"19", GitVersion:"v1.19.11", GitCommit:"c6a2f08fc4378c5381dd948d9ad9d1080e3e6b33", GitTreeState:"clean", BuildDate:"2021-05-13T17:14:18Z",
GoVersion:"go1.15.12", Compiler:"gc", Platform:"linux/amd64"}
```

Figure 10: kubectl version

3.4 Kubernetes Deployment:

A Kubernetes Deployment instructs Kubernetes on how to generate or change instances of pods that contain containerized applications. Eight YAML files make up the project, which is used to deploy everything in the Kubernetes cluster. Three of these files are deployed in the local machine, and the rest are configured for the Azure platform. The service name and the values of additional YAML files like configmap, secret, and seq log are all contained in the MVC deployment file. A load balancer is included in the file for a new service to make it accessible to others. A similar method is followed when deploying MySQL.

To handle individual components, it's time to put up an orchestrator like Kubernetes. Kubernetes contains several tools for scaling, networking, securing, and supporting containerized applications in addition to the capabilities of containers.

3.5 Describing apps using Kubernetes YAML

In Kubernetes, all containers are scheduled as pods, which are groupings of co-located containers and share some resources. Manifests called Kubernetes YAML files can and should be used to describe all Kubernetes objects. These YAML files explain all of a Kubernetes app's components and configurations and can be used to efficiently construct and destroy apps in any Kubernetes environment.

The configuration file for the web app in this project is called **cckubernetesproduct-mvc.deployment.azure.yaml**. For our mvc project, we've written deployment, service, and horizontal scaler pod code in this file. The default replica set in the deployment phase is 2. Our resources were similarly limited. We also specify that just port 80 is open. We get connection strings from secret for environment variables and build version and Seq port from the config map.

We create a service type loadbalancer on port 80 so that everyone may see our web app. The minimum pod count for horizontal pod scaling is 2, and the maximum pod count is 5. It will also work if the desired CPU utilization is more than or equal to 50.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: cckubernetesproduct-mvc
spec:
  selector:
    matchLabels:
      app: mvc
  replicas: 2 # tells deployment to run 2 pods matching the template
  template:
    metadata:
      labels:
        app: mvc
    spec:
      containers:
        - name: mvc
          image: aniksaha/cckubernetesproduct-mvc:v2.#{Build.BuildId}#
          ports:
            - containerPort: 80
          resources:
            limits:
              cpu: "0.4"
              memory: "200Mi"
            requests:
              cpu: "0.2"
              memory: "100Mi"
          env:
            - name: ConnectionStrings__ProductsContext
              valueFrom:
                secretKeyRef:
                  name: cckubernetesproduct-secret
                  key: db-connection-string-secret
            - name: ConnectionStrings__BuildVersion
              valueFrom:
                configMapKeyRef:
                  name: cckubernetesproduct-configmap
                  key: build-version
            - name: ConnectionStrings__Seq
              value: cckubernetesproduct-seq-log-service
            - name: ConnectionStrings__SeqPort
              valueFrom:
                configMapKeyRef:
                  name: cckubernetesproduct-configmap
                  key: seq-log-port
---
kind: Service
apiVersion: v1
metadata:
  name: cckubernetesproduct-mvc-service
spec:
  selector:
    app: mvc
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
---
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: cckubernetesproduct-mvc-hpa
spec:
  maxReplicas: 5 # define max replica count
  minReplicas: 2 # define min replica count
  scaleTargetRef:
    apiVersion: apps/v1
```

```

kind: Deployment
name: cckubernetesproduct-mvc
targetCPUUtilizationPercentage: 50 # target CPU utilization

```

In this file we have written deployment and service for MSSQL . In the deployment part, the default replica is set to 1. We here define one port 1433 that is accessible by others. For environment variables, we collect db password strings from Kubernetes secret.

This MSSQL is exposed to our web app only. So we make a service type Nodeport. That's how our db can not be accessible by others. So no one can temper our db data.

cckubernetesproduct-mssql.deployment.yaml

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: cckubernetesproduct-mssql
spec:
  replicas: 1
  selector:
    matchLabels:
      app: mssql
  template:
    metadata:
      labels:
        app: mssql
    spec:
      terminationGracePeriodSeconds: 10
      containers:
        - name: mssql
          image: microsoft/mssql-server-linux
          resources:
            limits:
              cpu: "1"
              memory: "1Gi"
            requests:
              cpu: "0.1"
          ports:
            - containerPort: 1433
          env:
            - name: ACCEPT_EULA
              value: "Y"
            - name: SA_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: cckubernetesproduct-secret
                  key: db-password-secret
          volumeMounts:
            - name: cckubernetesproduct-mssql-persistent-storage
              mountPath: /var/opt/mssql
      volumes:
        - name: cckubernetesproduct-mssql-persistent-storage
          persistentVolumeClaim:
            claimName: cckubernetesproduct-mssql-persistent-volume-claim
---
apiVersion: v1
kind: Service
metadata:
  name: cckubernetesproduct-mssql-service
spec:
  selector:
    app: mssql
  ports:
    - protocol: TCP
      port: 1433
      targetPort: 1433
      nodePort: 30200
  type: NodePort

```

In this file, we have written deployment and service for seq(logging web project). In the deployment part, we set the default replica as 1. We also limited our resources. We have to define two ports. One(5341) is for writing logging from a web project, and another port(80) is opened for showing it externally.

Here, two services as two ports are opened for this pod. Port 5341 is not opened for all. Only our web project will access it internally. So we define cckubernetesproduct-seq-log-service as no cluster IP. On the other hand, Port 80 can be accessed by all as we set type loadbalancer for cckubernetesproduct-seq-ui-service service.

cckubernetesproduct-seq.deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: cckubernetesproduct-seq
spec:
  selector:
    matchLabels:
      app: seq-app
  replicas: 1
  template:
    metadata:
      labels:
        app: seq-app
    spec:
      containers:
        - name: seq
          image: datalust/seq:latest
          ports:
            - containerPort: 5341
            - containerPort: 80
          resources:
            limits:
              cpu: "0.5"
              memory: "200Mi"
            requests:
              cpu: "0.25"
              memory: "100Mi"
          env:
            - name: ACCEPT_EULA
              value: "y"
---
apiVersion: v1
kind: Service
metadata:
  name: cckubernetesproduct-seq-ui-service
spec:
  selector:
    app: seq-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
---
apiVersion: v1
kind: Service
metadata:
  name: cckubernetesproduct-seq-log-service
spec:
  selector:
    app: seq-app
  ports:
    - protocol: TCP
      port: 5341
      targetPort: 5341
```

This YAML creates a ConfigMap for external configuration with the build version and the seq logport value that is set to 5341.

cckubernetesproduct-configmap.deployment.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: cckubernetesproduct-configmap
data:
  build-version: v2.#{Build.BuildId}#
  seq-log-port: "5341"
```

Here is a configuration file that was used to create secure credentials. All values inside this file are base64 encoded to not be stolen easily.

cckubernetesproduct-secret.deployment.yaml

```
apiVersion: v1
kind: Secret
metadata:
  name: cckubernetesproduct-secret
data:
  db-connection-string-
secret: U2VydmVpPWNja3ViZXJuc3Rlc3Byb2R1Y3QtbXNzcWwtc2Vydm1jZTtEYXRhYmFzZT1DQ0t1YmVybmV0ZXNQcm9kdWN0cztVc2VyPVNB01Bhc3N3b3JkPTEyMzQ1Njc4QWE7SW50ZWdyYXRlZCB7ZW50cm10eT1mYXNzZTtNdWw0aXBsZUFjdG12ZVJlc3VsdFNldHM9dHJ1ZTs=
  db-password-secret: MTIzNDU2NzhBYQ==
type: Opaque
```

To create Pods with persistent storage, we used a StorageClass and PersistentVolumeClaim. The configuration is similar to a Deployment but provides predictable names for the pods and allows us to add persistent storage classes. This yaml file defines persistent storage that can be mounted in containers. Here, this includes the ReadWriteOnce access mode so that the volume can be mounted by a single container in read-write mode. We also apply a size limitation to 1GB. The general structure of the configuration is as follows:

cckubernetesproduct-mssql-persistent-volume.deployment.azure.yaml

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: cckubernetesproduct-mssql-persistent-volume
provisioner: kubernetes.io/azure-disk
parameters:
  storageaccounttype: Standard_LRS
  kind: Managed
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: cckubernetesproduct-mssql-persistent-volume-claim
  annotations:
    volume.beta.kubernetes.io/storage-class: cckubernetesproduct-mssql-persistent-volume
spec:
  storageClassName: default
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
```

These are all configuration files used in this project. Now that we have the full definition of our application, all the particular servers can be deployed with the **kubectl apply** command specifying as a parameter the path of the YAML file which we have just created. For example, The following command is in use to deploy our web app,

-- kubectl apply cckubernetesproduct-mvc.deployment.azure.yaml

The Output shows the successfully created deployments and services.

3.6 Azure Kubernetes service:

AKS is one of many PaaS (Platform-as-a-Service) Azure services. It implies we won't have to worry about manually constructing, maintaining, and configuring virtual machines when we deploy a Kubernetes cluster [12].

We set up the Azure kubernetes service from the azure portal by selecting configuration from the web. For AKS it won't take that much time.

```
{
  "availabilityZones": [
    "1",
    "2",
    "3"
  ],
  "count": 2,
  "enableAutoScaling": false,
  "enableEncryptionAtHost": null,
  "enableFips": false,
  "enableNodePublicIp": null,
  "gpuInstanceProfile": null,
  "kubeletConfig": null,
  "kubeletDiskType": "OS",
  "linuxOsConfig": null,
  "maxCount": null,
  "maxPods": 110,
  "minCount": null,
  "mode": "System",
  "name": "agentpool",
  "nodeImageVersion": "AKSUbuntu-1804gen2c",
  "nodeLabels": {},
  "nodePublicIpPrefixId": null,
  "nodeTaints": null,
  "orchestratorVersion": "1.19.11",
  "osDiskSizeGb": 128,
  "osDiskType": "Managed",
  "osSku": "Ubuntu",
  "osType": "Linux",
  "podSubnetId": null,
  "powerState": {
    "code": "Running"
  },
  "provisioningState": "Succeeded",
  "proximityPlacementGroupId": null,
  "scaleSetEvictionPolicy": null,
  "scaleSetPriority": null,
  "spotMaxPrice": null,
  "tags": null,
  "type": "VirtualMachineScaleSets",
  "upgradeSettings": null,
  "vmSize": "Standard_B2s"
}
```

Figure 11: AKS Configuration

3.7 Kubernetes Dashboard

The Kubernetes Dashboard is a web-based UI for Kubernetes clusters that may be used for various purposes. It enables users to administer and debug cluster-based applications and operate the cluster itself [7].

To deploy Dashboard, the following command has executed:

```
--kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/v2.1.0/aio/deploy/recommended.yaml
```

The command to Create a Sample User Account that can Access the Dashboard via Token is:

```
--kubectl apply -f https://gist.githubusercontent.com/dahlsailrunner/bbd453f3bb6259b66c08a70d0908283f/raw/5727723217e2df4b65d8933adf04d009cfb0fe3f/local-dashboard-account.yaml
```

It is necessary to capture the Token. It's the token value returned by the command below.

PowerShell:

```
--kubectl -n kubernetes-dashboard describe secret $(kubectl -n kubernetes-dashboard get secret | sls admin-user |  
ForEach-Object { $_ -Split 's+' } | Select -First 1)
```

To create a secure route to the Kubernetes cluster from the local desktop to access Dashboard, the command below has executed:

```
--start kubectl proxy
```

3.8 Connect to AKS Dashboard and kubectl

The easiest way to connect to AKS cluster is to open the Azure dashboard and, in the Overview section, choose View Kubernetes dashboard [6]. From the right, a new panel with Azure CLI commands will appear.

To create the connection of the AKS cluster from a local machine, we follow some steps that set the --resource-group and the --name parameters with the specific values of our created cluster.

The following screenshot shows it:

```
PS C:\Users\Mitu Saha> az login
The default web browser has been opened at https://login.microsoftonline.com/common/oauth2/authorize. Please continue the login in the web browser. If no web browser is available or if the web browser fails to open, use device code flow with the command 'az login --use-device-code'.
You have logged in. Now let us find all the subscriptions to which you have access...
[
  {
    "cloudName": "AzureCloud",
    "homeTenantId": "979a61d1-be26-46a3-a669-c895b731071b",
    "id": "ca5bdee2-cbee-4572-ac5f-a2d1056e6394",
    "isDefault": true,
    "managedByTenants": [],
    "name": "Azure for Students",
    "state": "Enabled",
    "tenantId": "979a61d1-be26-46a3-a669-c895b731071b",
    "user": {
      "name": "shalim.sadman@northsouth.edu",
      "type": "user"
    }
  }
]
PS C:\Users\Mitu Saha> az account set --subscription ca5bdee2-cbee-4572-ac5f-a2d1056e6394
PS C:\Users\Mitu Saha> az aks get-credentials --resource-group CloudComputingGroup --name CCKuberneteters
Merged "CCKuberneteters" as current context in C:\Users\Mitu Saha\.kube\config
PS C:\Users\Mitu Saha>
```

Figure 12: Azure commands

The command below is precisely what we need to see on the dashboard. It will start a proxy on your local system that will launch the dashboard with data from the AKS instance.

--az aks browse --resource-group CloudComputingGroup --name CCKuberneteters

The dashboard provides information on the Kubernetes resources in cluster and any errors that may have occurred.

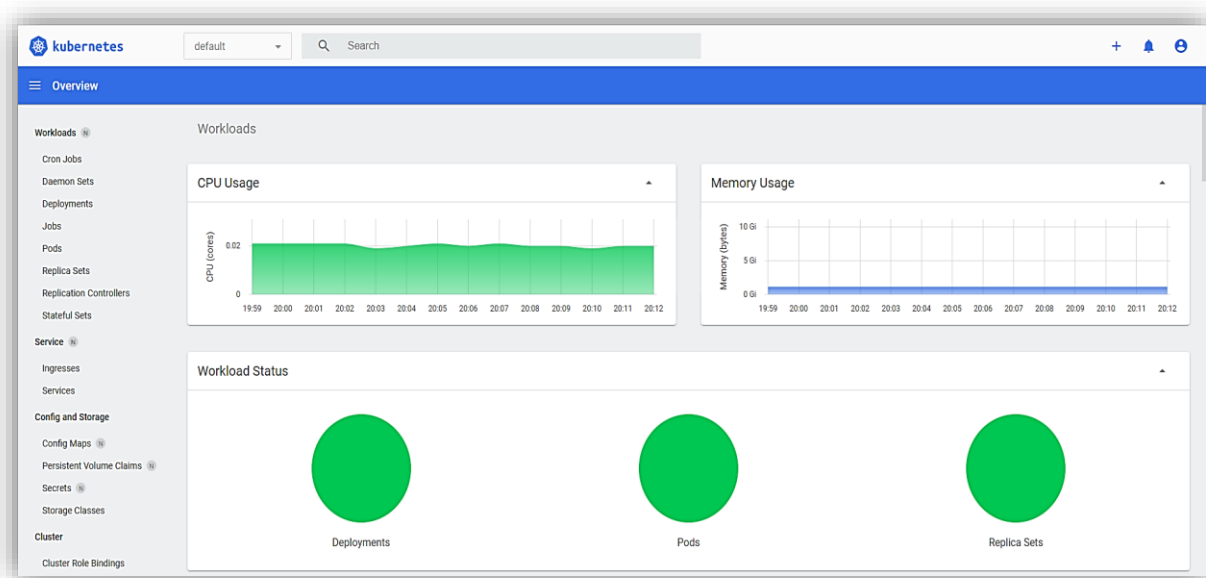


Figure 13: Kubernetes dashboard

The dashboard will review all the events that happened during the load test. In the Pods section, each pod shows the CPU spike it had:

Name	Namespace	Labels	Pods	Created ↑	Images	
✓ cckubernetesproduct-mssql	default	-	1 / 1	3 days ago	mcr.microsoft.com/mssql/server	⋮
✓ cckubernetesproduct-mvc	default	-	4 / 4	3 days ago	aniksaha/cckubernetesproduct-mvc:v4.14	⋮
✓ cckubernetesproduct-seq	default	-	1 / 1	3 days ago	datastus/seqlatest	⋮

1 ~ 3 of 3 |< < > >|

Name	Namespace	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Created ↑	
✓ cckubernetesproduct-mvc-6766bf794-gk4kp	default	app: mvc pod-template-hash: 6766bf794	aks-agentpool-91586894-vmss000000	Running	0			a minute ago	⋮
✓ cckubernetesproduct-mvc-6766bf794-z7cv8	default	app: mvc pod-template-hash: 6766bf794	aks-agentpool-91586894-vmss000001	Running	0			a minute ago	⋮
✓ cckubernetesproduct-mvc-6766bf794-lgtz8	default	app: mvc pod-template-hash: 6766bf794	aks-agentpool-91586894-vmss000000	Running	0			3 minutes ago	⋮
✓ cckubernetesproduct-mvc-6766bf794-djrkx	default	app: mvc pod-template-hash: 6766bf794	aks-agentpool-91586894-vmss000001	Running	0			4 minutes ago	⋮
✓ cckubernetesproduct-mssql-8486c69bdf-l27n	default	app: mssql pod-template-hash: 8486c69bdf	aks-agentpool-91586894-vmss000000	Running	0			3 days ago	⋮
✓ cckubernetesproduct-seq-667d99dfc6-cffk4	default	app: seq-app pod-template-hash: 667d99dfc6	aks-agentpool-91586894-vmss000000	Running	0			3 days ago	⋮

1 ~ 6 of 6 |< < > >|

Figure 14: deployment and pods

As a consequence, three new replica pods are running for the cckubernetesproduct. mvc deployment.

The following figure shows all the services and configmap that has been used:

Name	Namespace	Labels	Cluster IP	Internal Endpoints	External Endpoints	Created ↑	
✓ cckubernetesproduct-mvc-service	default	-	10.0.22.58	cckubernetesproduct-mvc-service:80 TCP cckubernetesproduct-mvc-service:30373 TCP	20.52.212.203:80	2 days ago	⋮
✓ cckubernetesproduct-mssql-service	default	-	10.0.36.193	cckubernetesproduct-mssql-service:1433 TCP cckubernetesproduct-mssql-service:30200 TCP	-	2 days ago	⋮
✓ cckubernetesproduct-seq-log-service	default	-	10.0.157.88	cckubernetesproduct-seq-log-service:5341 TCP cckubernetesproduct-seq-log-service:0 TCP	-	2 days ago	⋮
✓ cckubernetesproduct-seq-ui-service	default	-	10.0.49.102	cckubernetesproduct-seq-ui-service:80 TCP cckubernetesproduct-seq-ui-service:31657 TCP	20.79.72.53:80	2 days ago	⋮
✓ kubernetes	default	component: apiserver provider: kubernetes	10.0.0.1	kubernetes:443 TCP kubernetes:0 TCP	-	2 days ago	⋮

1 ~ 5 of 5 |< < > >|

Name	Namespace	Labels	Created ↑	
cckubernetesproduct-configmap	default	-	2 days ago	⋮

1 ~ 1 of 1 |< < > >|

Figure 15: Services, Config, and Storage

The subsequent figures show the AKS dashboard view of persistent volumes, secret, storage classes, and Replicas.

Persistent Volumes

Name	Capacity	Access Modes	Reclaim Policy	Status	Claim	Storage Class	Reason	Created ↑	
pvc-bd6518b4-516c-4992-ac57-affbf8fadc4e	storage: 1Gi	ReadWriteOnce	Delete	Bound	default/cckubernetes-mssql-persistent-volume-claim	cckubernetes-mssql-persistent-volume	-	2 days ago	
1 - 1 of 1									

Secrets

Name	Namespace	Labels	Type	Created ↑	
cckubernetesproduct-secret	default	-	Opaque	2 days ago	
default-token-65mkm	default	-	kubernetes.io/service-account-token	2 days ago	
1 - 2 of 2					

Storage Classes

Name	Provisioner	Parameters	Created ↑	
cckubernetes-mssql-persistent-volume	kubernetes.io/azure-disk	<div>kind: Managed</div> <div>storageaccounttype: Standard_LRS</div>	2 days ago	
azurefile	kubernetes.io/azure-file	<div>skuName: Standard_LRS</div>	2 days ago	
azurefile-premium	kubernetes.io/azure-file	<div>skuName: Premium_LRS</div>	2 days ago	
default	kubernetes.io/azure-disk	<div>cachingmode: ReadOnly</div> <div>kind: Managed</div> <div>storageaccounttype: StandardSSD_LRS</div>	2 days ago	

Figure 16: Persistent volumes, Secrets, Storage Classes















Replica Sets								
Name	Namespace	Labels	Pods	Created ↑	Images			
 cckubernetesproduct-mssql-68584678c5	default	app: mssql pod-template-hash: 68584678c5	0 / 0	2 days ago	mcr.microsoft.com/mssql/server			
 cckubernetesproduct-mssql-8486c69bdf	default	app: mssql pod-template-hash: 8486c69bdf	1 / 1	2 days ago	mcr.microsoft.com/mssql/server			
 cckubernetesproduct-mvc-5f5b447655	default	app: mvc pod-template-hash: 5f5b447655	2 / 2	2 days ago	aniksaha/cckubernetesproduct-mvc:v2.24			
 cckubernetesproduct-seq-667d99dfc6	default	app: seq-app pod-template-hash: 667d99dfc6	1 / 1	2 days ago	datalust/seq:latest			
1 ~ 4 of 4							   	

Figure 17: Replica Sets

3.9 AKS Cluster nodes

The following command provides us with pods, available services, deployments, and replica sets.

--kubectrl get all

```

PS C:\Users\Mitu Saha> kubectl get all
NAME                                     READY   STATUS    RESTARTS   AGE
pod/cckubernetesproduct-mssql-8486c69bdf-127tn   1/1     Running   0           3d20h
pod/cckubernetesproduct-mvc-6766bf794-djrkx       1/1     Running   0           3m35s
pod/cckubernetesproduct-mvc-6766bf794-gk4kp       1/1     Running   0           26s
pod/cckubernetesproduct-mvc-6766bf794-1gtz8       1/1     Running   0           3m28s
pod/cckubernetesproduct-mvc-6766bf794-z7cv8       1/1     Running   0           26s
pod/cckubernetesproduct-seq-667d99dfc6-cffk4      1/1     Running   0           3d21h

NAME                                     TYPE                      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
service/cckubernetesproduct-mssql-service        NodePort            10.0.36.193     <none>           1433:30200/TCP   3d20h
service/cckubernetesproduct-mvc-service          LoadBalancer       10.0.22.58      20.52.212.203    80:30373/TCP     3d20h
service/cckubernetesproduct-seq-log-service      ClusterIP           10.0.157.88     <none>           5341/TCP         3d21h
service/cckubernetesproduct-seq-ui-service       LoadBalancer       10.0.49.102     20.79.72.53      80:31657/TCP     3d21h
service/kubernetes                              ClusterIP           10.0.0.1        <none>           443/TCP          3d22h

NAME                                     READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/cckubernetesproduct-mssql      1/1     1             1           3d20h
deployment.apps/cckubernetesproduct-mvc        4/4     4             4           3d20h
deployment.apps/cckubernetesproduct-seq        1/1     1             1           3d21h

NAME                                     DESIRED   CURRENT   READY   AGE
replicaset.apps/cckubernetesproduct-mssql-68584678c5  0         0         0       3d20h
replicaset.apps/cckubernetesproduct-mssql-8486c69bdf  1         1         1       3d20h
replicaset.apps/cckubernetesproduct-mvc-5f5b447655   0         0         0       3d20h
replicaset.apps/cckubernetesproduct-mvc-6766bf794    4         4         4       3m36s
replicaset.apps/cckubernetesproduct-seq-667d99dfc6   1         1         1       3d21h

NAME                                     REFERENCE                                     TARGETS  MINPODS
horizontalpodautoscaler.autoscaling/cckubernetesproduct-mvc-hpa  Deployment/cckubernetesproduct-mvc  92%/50%  2

```

Figure 18: kubectl commands

Since the service type we have created is LoadBalancer, our application is exposed to the outside world through an external IP.

4. DevOps for Kubernetes using Azure DevOps

Microservices are growing increasingly popular in recent years. Kubernetes is where these microservices spend the majority of their time. A speedy and reliable deployment is a goal to achieve with microservices [11]. It will be demonstrated how to use Azure DevOps to set up a CI/CD pipeline to deploy the web app to a Kubernetes cluster.

4.1 Source control

We make a connection of our [GitHub project](#) with Azure DevOps.

4.2 Create a Build / Continuous Integration (CI) pipeline

The enabled Continuous integration lets the entire project work through a built-in Azure DevOps procedure (inside pipelines). Those steps include building and pushing images (in the docker hub), replacing build versions in yaml files, and copying and publishing artifacts at the end.

Here, for every new push in the Github repository, the continuous integration will fire up.

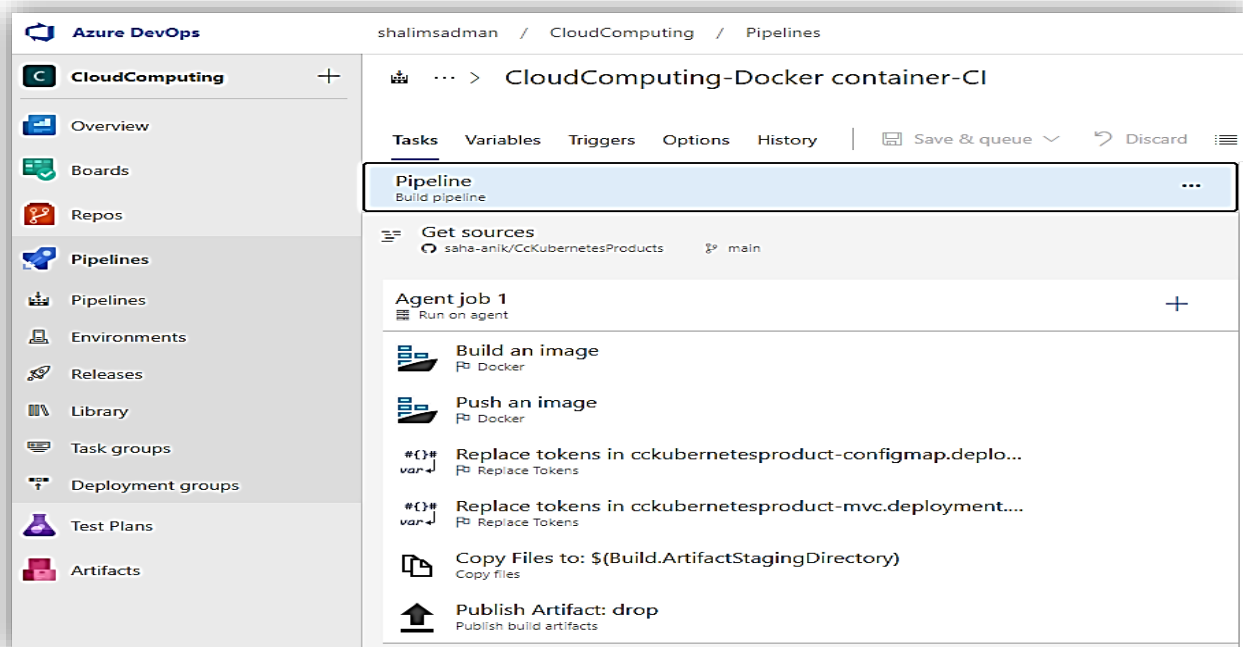


Figure 19: Continuous Integration (CI) pipeline

After running the web app pipeline, once the build process starts, the following build jobs occur in progress.

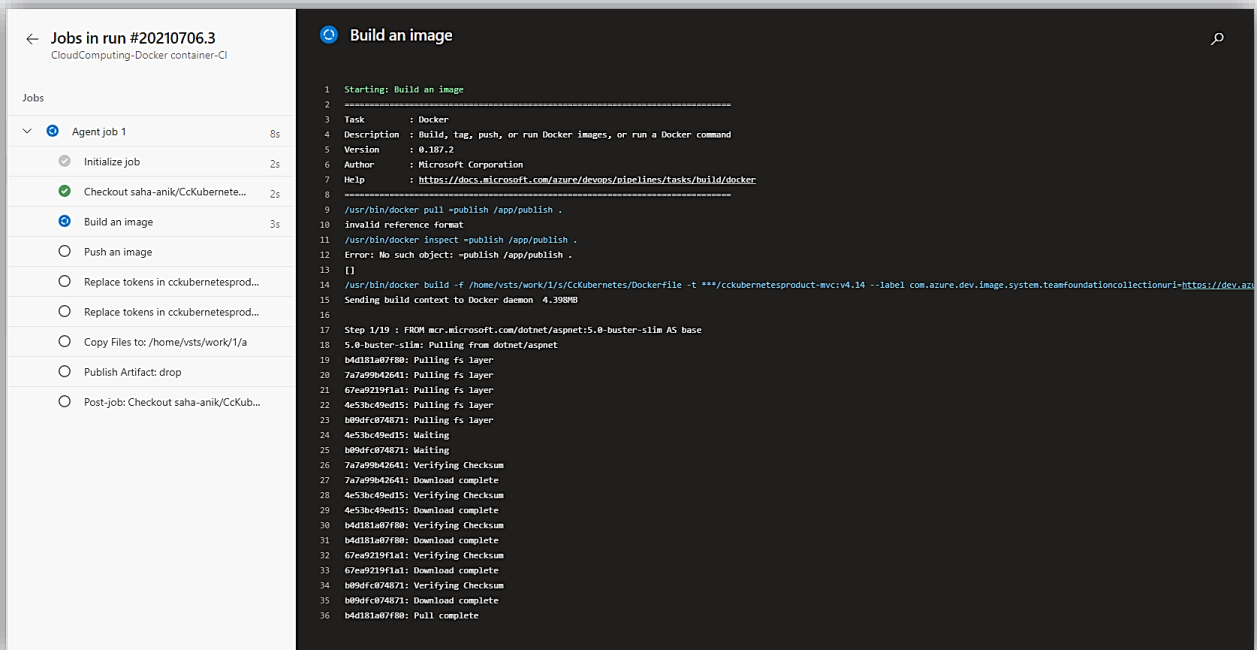


Figure 20: Progress jobs in Continuous Integration

These are the brief descriptions of all executed pipelines in DevOps.

All pipeline runs		
Description	Stages	
add script for load testing #20210623.2 on cckubernetesproduct main b01f6b4	✓	Jun 23 1m 24s
add hpa #20210623.1 on cckubernetesproduct main 0991aba	✓	Jun 23 1m 41s
add string for seq log #20210622.14 on cckubernetesproduct main 430b2e6	✓	Jun 22 1m 17s
fix port issue for seq log #20210622.13 on cckubernetesproduct main 3ebab58	✓	Jun 22 1m 56s
cluster ip #20210622.12 on cckubernetesproduct main dd53e5c	✓	Jun 22 1m 27s
change azure #20210622.11 on cckubernetesproduct main 3c7f594	✓	Jun 22 1m 28s
change index #20210622.10 on cckubernetesproduct main 0a34f0b	✓	Jun 22 1m 19s

Figure 21: Executed pipeline in Continuous Integration

4.3 Create a Release / Continuous Delivery (CD) pipeline

In the release, we apply kubectl to configure and connect all the YAML files of k8s components with azure k8s. An Azure release pipeline is created for the mvc app to be able to deploy it via kubectl. The **Releases** tab in the **Pipelines** section shows each command that connects azure pipelines to Kubernetes.

The screenshot displays the Azure DevOps web interface. On the left, a sidebar contains navigation links: Overview, Boards, Repos, Pipelines (selected), Environments, Releases, Library, Task groups, Deployment groups, Test Plans, and Artifacts. The main area shows the 'New release pipeline' configuration page. At the top, there's a breadcrumb trail: 'All pipelines > New release pipeline'. Below this, tabs for 'Pipeline', 'Tasks', 'Variables', 'Retention', 'Options', and 'History' are visible, with 'Tasks' being the active tab. A 'Prod' deployment process is selected. Under the 'Agent job' section, a list of tasks is shown, each with a Kubernetes icon and the command 'kubectl apply -secret', 'kubectl apply - configmap', 'kubectl apply -persistant volume', 'kubectl apply -mssql', 'kubectl apply - seq', and 'kubectl apply -mvc'.

Figure 22: Release pipelines

The continuous delivery process will start after each successful continuous integration.

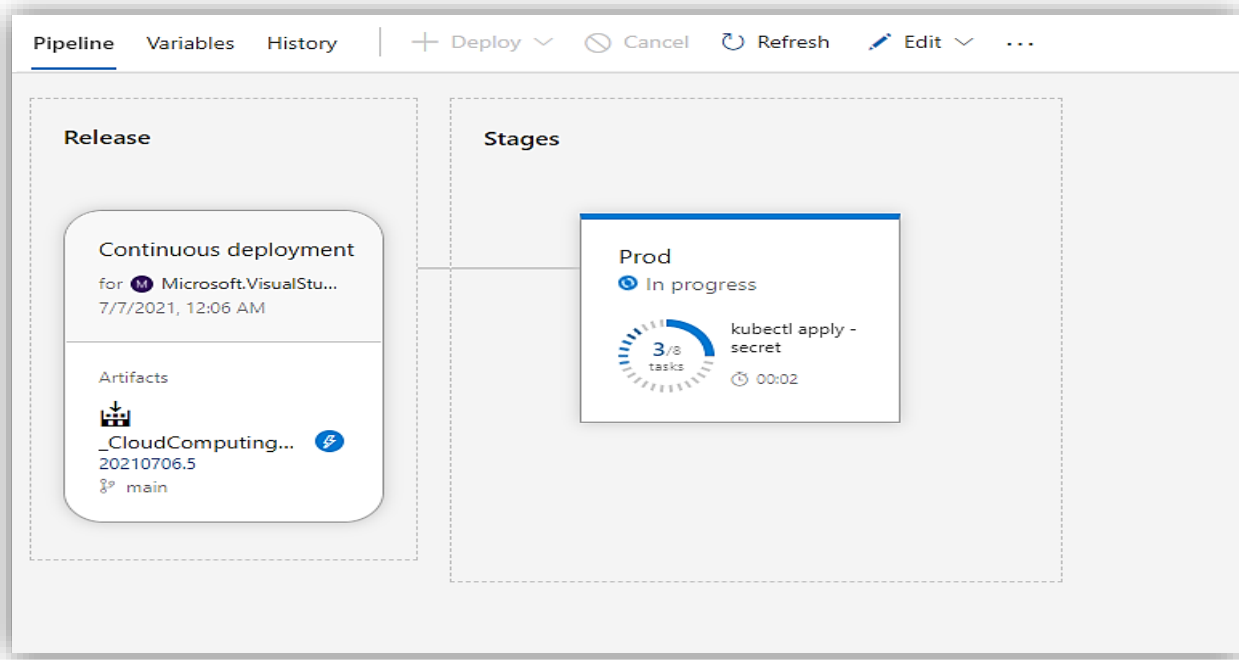


Figure 23: Continuous delivery

Finally, the web app has deployed along with MySQL and seq in a Kubernetes cluster by leveraging Azure DevOps, Azure Kubernetes Service, Docker.

5. Load Testing:

Load testing evaluates a software application's performance under a specific load. It determines how the software application operates when multiple users access it at the same time. Load testing is used to identify performance bottlenecks and ensure software applications' stability and smooth operation before deployment [13].

To test horizontal pod autoscaling, we use load testing. For our account online, we have two pods. We increase the number of pods when CPU utilization reaches 50%. Our web project has a maximum of 5 pods. We utilize the k6 framework for load testing because it is an open-source load testing tool. The k6 APIs are simple to use, versatile, and capable. Writing tests in JavaScript is similarly simple [14].

```

PS D:\his\Cloud Computing\CloudComputing\Cckubernetes> k6 run .\loadtesting.js

      /\_/\
     /  _  \
    /_____\ \
   /         \
  /           \
 /             \
/               \
\               /
 \             /
  \           /
   \         /
    \_____/

execution: local
script: .\loadtesting.js
output: -

scenarios: (100.00%) 1 scenario, 500 max VUs, 2m10s max duration (incl. graceful stop):
 * default: 500 looping VUs for 1m40s (gracefulStop: 30s)

running (1m41.0s), 000/500 VUs, 48231 complete and 0 interrupted iterations
default ✓ [=====] 500 VUs 1m40s

data_sent.....: 3.8 MB 38 kB/s
http_req_blocked.....: avg=5.26ms min=0s med=0s max=1.08s p(90)=0s p(95)=0s
http_req_connecting.....: avg=5.25ms min=0s med=0s max=1.08s p(90)=0s p(95)=0s
http_req_duration.....: avg=32.41ms min=6.98ms med=24.09ms max=597.99ms p(90)=44.13ms p(95)=64.02ms
{ expected_response:true }...: avg=32.41ms min=6.98ms med=24.09ms max=597.99ms p(90)=44.13ms p(95)=64.02ms
http_req_failed.....: 0.00% ✓ 0 X 48231
http_req_receiving.....: avg=296.65µs min=0s med=0s max=109.51ms p(90)=826.5µs p(95)=1ms
http_req_sending.....: avg=28.41µs min=0s med=0s max=7.07ms p(90)=0s p(95)=152.15µs
http_req_tls_handshaking.....: avg=0s min=0s med=0s max=0s p(90)=0s p(95)=0s
http_req_waiting.....: avg=32.09ms min=6.98ms med=23.85ms max=597.28ms p(90)=43.74ms p(95)=63.15ms
http_reqs.....: 48231 477.41459/s
iteration_duration.....: avg=1.04s min=1s med=1.02s max=2.67s p(90)=1.04s p(95)=1.07s
iterations.....: 48231 477.41459/s
vus.....: 37 min=37 max=500
vus_max.....: 500 min=500 max=500

```

Figure 24: Iterations after executing load testing

our load testing js script is as follows:

```

import http from 'k6/http';
import { sleep } from 'k6';
export let options = {
  vus: 200,
  duration: '60s',
};
export default function () {
  http.get('http://20.79.117.111/Products');
  sleep(1);
}

```

6. Conclusion & future work:

With the introduction of the cloud, the computing world has forever changed. Cloud computing provides developers with rapid, low-cost access to infrastructure at nearly unlimited scales. Because of the cloud's agility and high availability, monolithic architectures have been stressed, resulting in the growth of microservices-based systems. In this paper, we've covered the fundamentals of microservices, Kubernetes, and Azure DevOps, as well as how to use them to deploy a simple web application. It contains detailed instructions for configuring and installing the required tools and services. Finally, we went over how to get to the AKS cluster's Kubernetes web interface, test deployments, and represent the deployment results. We also integrated the CI/CD Pipeline with Azure DevOps. To ensure that our deployment was stable, we performed load and log testing.

With all of the items that must be handled, such as ConfigMaps, services, pods, and Persistent Volumes, as well as the number of releases that must be maintained, the system we worked on might be quite complex. As a result, we'd like to use Helm to parameterize our YAML in the future. It's a package manager that makes application deployment simple, standardized, and reusable, boosting developer productivity, reducing deployment complexity, improving operational readiness, and accelerating the adoption of cloud-native apps.

7. Reference:

1. Kubernetes (K8s). URL: <https://github.com/kubernetes/kubernetes> (visited on 27/06/2021).
2. WHAT IS KUBERNETES? URL: <https://kemptechnologies.com/blog/what-is-kubernetes/?fbclid=IwAR07gKxCVYy1LgYqJvTfPMhYy2tZHLniHi625vyWi9T7xLZYsW3n28sbiY4g> (visited on 08/07/2021)
3. Azure Kubernetes Service (AKS). URL: <https://azure.microsoft.com/en-us/services/kubernetes-service/#features> (visited on 01/07/2021).
4. Developing Microservices Architecture on Microsoft Azure with Open Source Technologies. URL: https://www.microsoftpressstore.com/store/developing-microservices-architecture-on-microsoft-9780136819387?fbclid=IwAR0eC35v3h0_QnH9bQ-Uf2guYMZ7xHryjoRbk1VX-6haaR-RoBvEuMel-sA (visited on 30/06/2021).
5. Kubernetes vs. Docker. URL: https://azure.microsoft.com/en-us/topic/kubernetes-vs-docker/?fbclid=IwAR2V4Eyvs6bl4DdkoqXYdZHKFxOGmrQUSUeWIabw4hyfewujUF_RNvPFVE (visited on 01/07/2021).
6. First steps with Docker and Kubernetes - Moving to Azure Kubernetes Service. URL: <https://techcommunity.microsoft.com/t5/windows-dev-appconsult/first-steps-with-docker-and-kubernetes-moving-to-azure/ba-p/360010> (visited on 02/07/2021).
7. Using the K8s Dashboard Locally. URL: <https://gist.github.com/dahlsailrunner/1a47b0e38f6e3ba64d4d61835c73b7e2?fbclid=IwAR1d-R9HCgnnL9By9WHDzCbUKhXHhcFWFN-Z7OZhS-cbelXWP6vwXrBEjZc> (visited on 03/07/2021).
8. How-To deploy Docker images to Azure Kubernetes Services (AKS). URL: <https://purple.telstra.com.au/blog/how-to-deploy-docker-images-to-azure-kubernetes-services-aks> (visited on 02/07/2021).
9. SEQ, Machine data, for humans. URL: <https://datalust.co/seq> (visited on 02/07/2021).
10. Kubernetes vs Docker in 2020. URL: https://technofaq.org/posts/2020/06/kubernetes-vs-docker-in-2020/?fbclid=IwAR11bbEILMFrCA_PDCvJ4HJQEjz1i5C5lGkLUxZ-FOHh7SC33l65yRxqG3Q (Visited on 08/07/2021)
11. Using Azure DevOps to setup a CI/CD pipeline and deploy to Kubernetes. URL: <https://cloudblogs.microsoft.com/opensource/2018/11/27/tutorial-azure-devops-setup-cicd-pipeline-kubernetes-docker-helm/> (visited on 01/07/2021).
12. Deploying a multi-container application to Azure Kubernetes Services. URL: [https://azuredevopslabs.com/labs/vstsextend/kubernetes/#:~:text=Azure%20Kubernetes%20Service%20\(AKS\)%20is,applications%20without%20container%20orchestration%20expertise](https://azuredevopslabs.com/labs/vstsextend/kubernetes/#:~:text=Azure%20Kubernetes%20Service%20(AKS)%20is,applications%20without%20container%20orchestration%20expertise) (visited on 28/06/2021).
13. Load Testing Tutorial: What is? How to? URL: <https://www.guru99.com/load-testing-tutorial.html> (visited on 02/07/2021).
14. The best developer experience for load testing. URL: <https://k6.io/> (visited on 02/07/2021).
15. Microservices with Azure Kubernetes and Docker. URL: <https://medium.com/@sumindaniro/microservices-with-azure-kubernetes-and-docker-49de617f0341> (visited on 04/07/2021).
16. Introduction to Docker and Kubernetes on Azure. URL: <https://dzone.com/articles/introduction-docker-and> (visited on 03/07/2021).
17. Build and deploy a microservice with Kubernetes. URL: https://searchitoperations-techtargt.com.cdn.ampproject.org/v/s/searchitoperations.techtargt.com/tutorial/Build-and-deploy-a-microservice-with-Kubernetes?amp_js_v=a6&_gsa=1&=1&usqp=mq33lAQKKAfQArABIIACAw%3D%3D#aoh=16256498080184&_ct=1625650078960&referrer=https%3A%2F%2Fwww.google.com&_tf=From%20%251%24s&share=https%3A%2F%2Fsearchitoperations.techtargt.com%2Ftutorial%2FBuild-and-deploy-a-microservice-with-Kubernetes (visited on 04/07/2021).